Deep Learning-Based Medication Recommendation System for Primary Care Decision Support

Arnav Sahu*, Hrithiq Gupta*, Padavinangady Nakul Bhat*, Pratyush Bhagat*
*Department of Computer Science and Engineering
Manipal Institute of Technology, Manipal, India
Email: arnav5.mitmpl2023@learner.manipal.edu
hrithiq.mitmpl2023@learner.manipal.edu
nakul1.mitmpl2023@learner.manipal.edu
pratyush3.mitmpl2023@learner.manipal.edu

Abstract

Medication recommendation in outpatient environments is a multi-label decision problem that must balance efficacy, patient-specific constraints and drug–drug safety. This paper (1) presents a detailed, paraphrased and survey of twenty important studies in representation learning, temporal modeling, graph-based drug recommendation and interpretability; (2) synthesizes findings and gaps; and (3) proposes a hybrid outpatient-focused deep learning architecture with integrated safety filtering and local explanations.

Index Terms

Medication recommendation, clinical decision support, graph neural networks, transformer, interpretability.

I. Introduction

Modern medication prescribing is a complex clinical decision-making process that requires balancing therapeutic efficacy, patient-specific constraints, dosage considerations, and potential adverse drug interactions. Physicians must select from thousands of available drug options while accounting for comorbidities, demographics, and safety risks. This cognitive burden often leads to variability in prescribing patterns and increases the risk of suboptimal or unsafe medication combinations. As electronic health records (EHRs) have become widely available, researchers have increasingly explored data-driven approaches to assist in medication recommendation and clinical decision support.

Deep learning (DL) methods have demonstrated strong capability in modeling complex, high-dimensional healthcare data. Early work such as *Doctor AI* [1] employed recurrent neural networks (RNNs) over longitudinal EHR data to predict future diagnoses and medication categories. By learning temporal patient embeddings from sequences of visits, these models captured disease progression patterns and improved multi-label prediction performance. Similarly, *RETAIN* [2] introduced a reverse-time attention mechanism that enhanced interpretability while maintaining strong predictive accuracy. By assigning attention weights to past visits and clinical variables, RETAIN allowed clinicians to better understand model reasoning, addressing a key concern in deploying deep learning systems in healthcare.

Beyond sequence modeling, recent research has incorporated graph-based representations and external medical knowledge to improve safety and robustness. *GAMENet* [3] integrated patient history encoding with a Graph Convolutional Network (GCN) operating over drug–drug interaction (DDI) graphs. By explicitly modeling adverse drug interactions, GAMENet reduced unsafe medication combinations while preserving recommendation accuracy. Other approaches, such as *SafeDrug* [4], leveraged molecular graph encoders to capture chemical substructures and improve DDI-aware predictions. These advances illustrate how deep architectures can combine temporal modeling, graph neural networks, and domain knowledge to enhance both effectiveness and safety in medication recommendation systems.

However, most existing approaches assume rich, longitudinal EHR data typically available in tertiary-care or inpatient hospital settings. In contrast, outpatient and primary-care environments often involve sparse data, limited historical visits, and short free-text symptom descriptions. Many encounters consist of a single visit with minimal structured history, which reduces the effectiveness of purely sequential models. Furthermore, outpatient medication guidance must explicitly prioritize safety, ensuring that predicted drug combinations minimize harmful interactions even when historical context is limited.

In this work, we address these challenges by conducting a focused survey of deep learning-based medication recommendation systems, with particular emphasis on safety modeling, interpretability, and applicability to data-sparse outpatient settings. Based on identified gaps, we propose a hybrid deep learning architecture that integrates attention-based symptom encoding, structured feature fusion, and DDI-aware safety filtering. Our objective is to design a medication recommendation framework that balances predictive performance, safety constraints, and interpretability, making it suitable for primary-care clinical decision support.

II. Literature Review

Recent research on medication recommendation systems has increasingly leveraged deep learning to address the complexity, scale, and safety requirements of clinical decision-making. Early approaches focused on sequential modeling of electronic health records (EHRs) using recurrent neural networks to capture temporal disease progression and prescribing patterns. More recent developments have incorporated attention mechanisms, graph neural networks, and knowledge-based constraints to improve interpretability and reduce adverse drug–drug interactions. In parallel, transformer-based language models and transfer learning techniques have shown strong performance in extracting clinically meaningful representations from short, unstructured symptom text. Together, these trends reflect a shift toward hybrid architectures that combine temporal modeling, structured medical knowledge, and explainable AI to produce safer and more reliable medication recommendations, particularly in data-sparse healthcare settings.

1. **DP-CRNN: CNN–RNN for pre-diagnosis (Zhou et al.)**
   Zhou et al. propose a double-pooling CRNN that first extracts n-gram semantics via convolutional filters, then models sequential patterns with recurrent layers and applies a second pooling step to emphasize salient temporal features. The architecture is well-suited to consumer-style inputs but does not extend to multi-drug prescription directly.
2. **Deep Patient (Miotto et al.)**
   Miotto et al. train stacked denoising autoencoders on large EHR collections to produce latent patient vectors usable in downstream prediction tasks. It demonstrates that unsupervised representation learning captures clinically meaningful structure, though it lacks temporal ordering modeling.
3. **Doctor AI (Choi et al.)**
   Doctor AI uses recurrent neural networks on sequences of visits to forecast diagnoses and medications. It highlights the importance of sequential modeling, but interpretability is limited and it does not account explicitly for drug–drug interactions.
4. **RETAIN (Choi et al.)**
   Introduces a reverse-time attention mechanism that affords clinicians interpretable importance scores for historical events. It offers improved transparency compared to black-box RNNs but was not designed for multi-drug safety constraints.
5. **GAMENet (Shang et al.)**
   Combines a memory-augmented patient encoder with graph-augmented drug representations to recommend medication sets while reducing harmful combinations. Its dependency on rich EHR data limits direct applicability to sparse outpatient inputs.
6. **SafeDrug (Lee et al.)**
   Employs molecular graph encoders to capture chemical substructures and merges them with clinical context. By modeling molecular features, it reduces adverse interactions but increases computational overhead.
7. **GRAM (Choi et al.)**
   Enhances representation by injecting hierarchical medical ontology structure (e.g., ICD tree) via graph attention, improving performance on rare codes.
8. **LEAP (Zhang et al.)**
   Frames multi-drug prescription as an optimization balancing efficacy and safety. It models combinations directly but requires long-term outcome labels rarely available in outpatient datasets.
9. **Molecular/Substructure models (MoleRec)**
   Molecular-focused methods use message-passing networks to learn drug representations from chemical graphs. They aid in predicting interactions but add model complexity.
10. **ClinicalBERT / Transformer-based models**
    Clinical adaptations of BERT produce strong contextual embeddings for clinical notes. These excel at capturing symptom semantics but are computationally heavy.
11. **Memory-Augmented Networks**
    These store past visit vectors to inform current recommendations. These models assume sufficient patient history, which is often missing in outpatient use-cases.
12. **Hierarchical Attention (HAN)**
    Processes multi-granular text (words $\rightarrow$ sentences) and highlights relevant passages for interpretability, beneficial for clinician-facing explanations.
13. **GRU-D (Che et al.)**
    Incorporates missingness masking and decay mechanisms tailored for clinical time series with irregular sampling, improving robustness for lab/vital-based predictions.
14. **Polypharmacy Side-Effect Modeling (Zitnik et al.)**
    Models adverse outcomes of drug combinations using graph neural approaches on drug–drug networks, enabling direct

estimation of interaction probabilities.

15. **Transfer Learning for Clinical NLP**
    Pretrained models show gains when labeled data are scarce, enabling better performance for symptom-to-disease mapping from small datasets.
16. **Multi-Task and Multi-Label Models**
    Jointly predict diagnoses, urgency, and medication sets, exploiting shared structure. Proper task weighting is crucial to avoid negative transfer.
17. **Reinforcement Learning for Treatment Policies**
    Optimizes sequential treatment decisions for long-term outcomes. Requires longitudinal reward signals and carries safety concerns if deployed naively.
18. **Explainable AI (LIME/SHAP)**
    Model-agnostic local explanation tools provide instance-level rationales; integrating graph- or rule-based explanations yields stronger clinical trust.
19. **Knowledge Graph Embedding Approaches**
    Connects diseases, drugs, and procedures in a structured latent space. They require curation and harmonization across various sources.
20. **Large Clinical Transformer Architectures**
    Trained on mixed clinical signals (notes + structured), these achieve state-of-the-art performance but demand significant compute and privacy considerations.

In summary, the surveyed literature demonstrates significant progress in applying deep learning to medication recommendation, spanning sequential models, attention-based architectures, graph neural networks, molecular representations, and large-scale transformers. While many approaches achieve strong predictive performance, most rely on dense longitudinal EHR data and are tailored to inpatient or ICU settings. Safety-aware models incorporating drug–drug interaction knowledge have emerged, yet often at the cost of increased architectural complexity or data requirements. Moreover, interpretability and outpatient generalization remain open challenges. These limitations highlight the need for lightweight, safety-aware, and explainable deep learning frameworks that can operate effectively under sparse data conditions typical of primary-care environments.

From the above analyses we identify persistent gaps:

- **Outpatient generalization:** Most safety-aware methods assume lengthy EHR histories; outpatient inputs are short and sparse.
- **Integrated safety + explainability:** Few systems simultaneously provide DDI-aware decisions and human-interpretable rationales.
- **Dosage and demographic adaptation:** Most research suggests drugs but not patient-specific dose ranges adjusted for weight/age/kidney function.

TABLE I
KEY PROPERTIES OF THE TWENTY SURVEYED STUDIES (PARAPHRASED).

#	Study / Year	Architecture	Primary Data	Safety Modeling	Main Limitation
1	DP-CRNN (Zhou et al., 2021)	CNN + RNN (double-pooling)	Short symptom text	No	Not drug-centric
2	Deep Patient (Miotto et al., 2016)	Stacked autoencoders	EHR (structured)	No	Static representations
3	Doctor AI (Choi et al., 2016)	RNN	Visit sequences	No	Low interpretability
4	RETAIN (Choi et al., 2016)	RNN + reverse-time attention	EHR seq.	No	Not drug-aware
5	GAMENet (Shang et al., 2019)	GNN + Memory	EHR + DDI graph	Yes (DDI-aware)	Needs dense EHR
6	SafeDrug (Lee et al., 2021)	Molecular GNNs	Drug graphs + EHR	Yes (molecular)	Compute-heavy
7	GRAM (Choi et al., 2017)	GAT over ontology	EHR + ontology	Partial	Ontology-dependency
8	LEAP (Zhang et al., 2017)	Optimization + Seq models	EHR	Partial	Requires outcome labels
9	MoleRec (various, 2022–23)	Message-passing GNNs	Molecular substructures	Yes	Complex training
10	ClinicalBERT (Alsentzer et al., 2019)	Transformer (BERT)	Clinical notes	No	Pretrain cost
11	MemoryNet (Wang et al., 2019)	Memory-augmented NN	Visit history	No	Needs history
12	HAN (Yang et al., 2016)	Hierarchical attention	Clinical text	No	Complex
13	GRU-D (Che et al., 2018)	GRU with decay	Irregular time-series	No	Focused on vitals/labs
14	Polypharmacy GNN (Zitnik et al., 2018)	GNN	Drug–drug network	Yes	Graph quality sensitive
15	Transfer clinical NLP (Lee et al., 2020)	Pretrained embeddings	Text	No	Domain shift issues
16	Multi-task CDSS (Harutyunyan et al., 2019)	Multi-task DNN	ICU time-series	Partial	ICU-focused
17	RL for treatments (Komorowski et al., 2018)	RL policies	Longitudinal EHR	Partial	Safety / reward design
18	LIME/SHAP explainability (Ribeiro/Lundberg)	Model-agnostic explainers	Any model	No (explain only)	Approximate explanations
19	Knowledge Graph Embeddings (various)	KG embeddings	KG (disease/drug)	Yes (KG rules)	KG curation needed
20	Large Clinical Transformers (2020–22)	Transformers	Mixed (notes+EHR)	Partial	Resource intensive

III. Methodology

The proposed methodology follows a pipeline-oriented approach designed to transform sparse, unstructured primary-care data into safe and explainable medication recommendations. The architecture is divided into four main phases: Data Acquisition and Preprocessing, Patient Representation Learning, Recommendation Logic and Safety Filtering, and Interpretability Mapping.

A. Data Acquisition and Preprocessing

The system first ingests clinical inputs which, in an outpatient setting, primarily consist of short-text symptom descriptions and basic structured attributes (age, sex, and known comorbidities).

- **Text Normalization:** To handle the noise in primary-care notes, we apply spell-correction and expand medical abbreviations (e.g., "SOB" to "Shortness of Breath").
- **Medical Entity Extraction:** Symptoms are mapped to standardized medical ontologies (e.g., SNOMED-CT or ICD-10) using a Named Entity Recognition (NER) layer to ensure semantic consistency.

B. Patient Representation Learning

We utilize a multi-modal embedding strategy to capture the full context of a patient encounter:

- **Symptom Encoding (E_s):** A lightweight Transformer-based encoder (Clinical DistilBERT) processes the normalized text to generate a dense semantic vector. This captures the nuance of patient complaints without the high latency of larger models.
- **Feature Fusion:** Structured demographic data and categorical comorbidity flags are passed through a multi-layer perceptron (MLP) and concatenated with the symptom vector E_s. The resulting joint embedding, P_v, represents the comprehensive patient state.

C. Recommendation Logic and Safety Layer

The recommendation process is framed as a multi-label classification problem, where the model predicts the probability of N possible drug classes.

- **Multi-Label Head:** The fusion vector P_v is passed through a Sigmoid output layer. To account for dependencies between medications (e.g., certain drugs are frequently prescribed together), we apply a label-correlation regularizer during training.
- **DDI-Aware Safety Filter:** Predicted candidates are passed through a graph-based safety filter. We construct a Drug-Drug Interaction (DDI) Graph using datasets like DrugBank. If the model suggests a combination with a high-risk interaction score, the safety layer either down-ranks the risky medication or flags it for clinician override.

D. Interpretability and Feedback

To ensure clinical trust, the system generates dual-layered explanations:

- **Saliency Mapping:** Using SHAP (SHapley Additive exPlanations), the model highlights specific words in the input symptoms that most influenced the recommendation.
- **Safety Rationales:** For any flagged interactions, the system provides a knowledge-graph trace (e.g., "Drug A and Drug B both inhibit enzyme CYP3A4"), offering a transparent reason for the safety alert.

E. Medication Recommendation Algorithm

The workflow of the proposed deep learning based medication recommendation system is summarized in Algorithm 1. The algorithm processes patient symptom text and structured demographic features, learns a patient representation using neural encoders, and generates medication recommendations while ensuring safety through a drug–drug interaction filter.

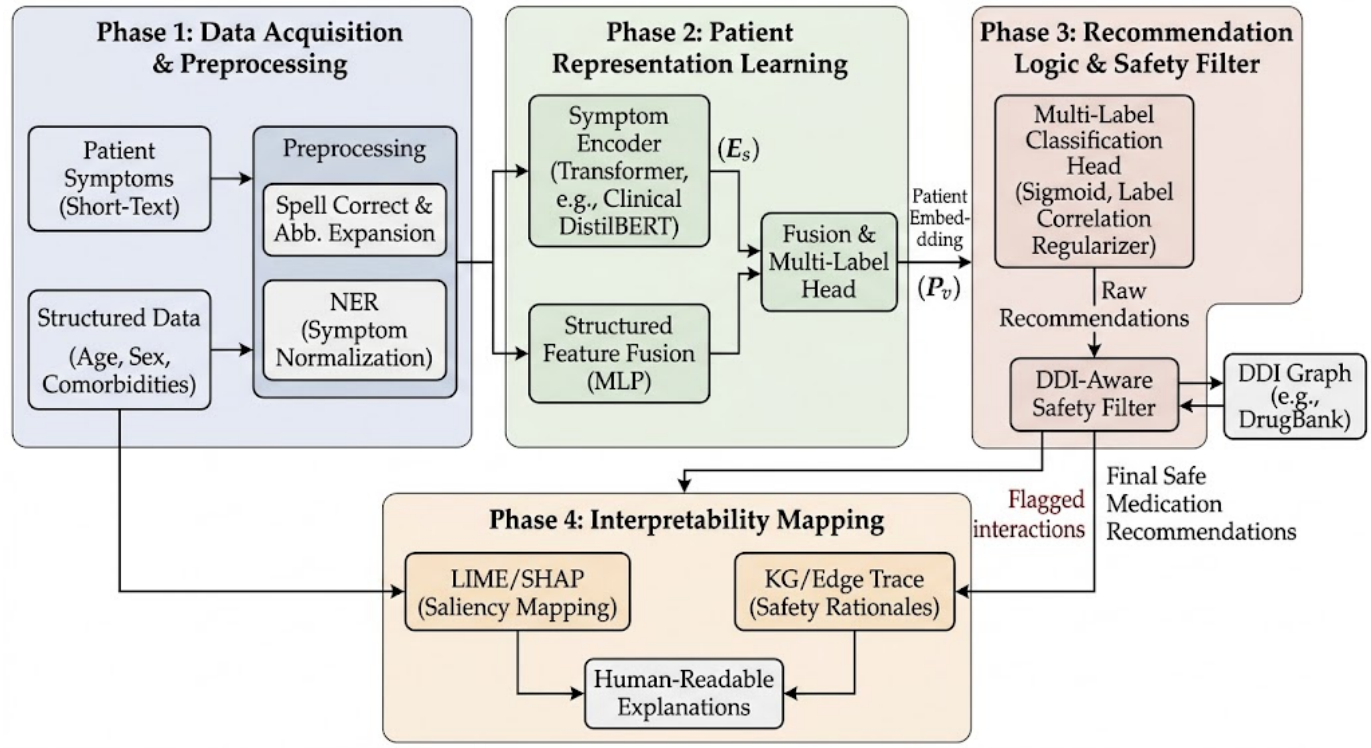

Fig. 1. Architecture Diagram detailing the four phases of the proposed medication recommendation methodology, including data flow, key modules, and the integrated safety and interpretability layers.

Algorithm 1 DL_Medication_Recommendation_System

Require: Patient symptom text S, structured data D (age, sex, comorbidities), DDI graph G
Ensure: Safe medication recommendations M_{safe} and explanation report E

Collect patient symptoms S and structured attributes D
Preprocess symptoms using spell correction and abbreviation expansion
Extract medical entities using Named Entity Recognition (NER)
Convert extracted symptoms into standardized medical terms
Generate symptom embedding E_s using a transformer encoder
Encode structured features D using a multi-layer perceptron
Fuse embeddings to obtain patient representation

$$P_v = Concatenate(E_s, E_d)$$

Predict medication probabilities using a neural classifier
Select top-k medications as candidate recommendations
Apply drug–drug interaction filtering using graph G
Generate explanation scores for model predictions
return safe medications M_{safe} and explanation report E

IV. Evaluation Plan

- **Data:** Synthea synthetic outpatient encounters, augmented symptom–disease datasets, curated DrugBank/TWOSIDES for safety.
- **Metrics:** Disease-prediction (accuracy, macro-F1), medication recommendation (Jaccard index, MAP), DDI rate, clinician-rated explanation utility.
- **Ablation:** Encoder ablation, safety-layer ablation, explanation fidelity tests.

V. Experimental Setup

A. Dataset

Due to the absence of a single publicly available dataset containing all required clinical attributes, we constructed a synthetic yet clinically grounded dataset by integrating multiple publicly available sources. Symptom–disease mappings were sourced from a structured `symptom_disease.csv` file, drug–disease associations from `drug_disease.csv`, and known adverse drug pair interactions from `ddi.csv`. Demographic attributes such as age and gender were drawn from Synthea-generated synthetic patient records (`patients.csv`), with clinical encounter data used to inform realistic naming conventions and distributions.

A custom data generation pipeline was implemented to simulate realistic outpatient encounters. Each sample was constructed by randomly selecting one to two diseases from the disease vocabulary, sampling associated symptoms with variable-length lists and applying text augmentation techniques including synonym replacement (e.g., *fever* $\rightarrow$ *pyrexia*) and template-based sentence construction (e.g., "patient reports ...", "complains of ..."). Medications were assigned via disease–drug mappings with fallback to global drug pool sampling. Drug combinations were filtered using the DDI dataset to remove pairs with known adverse interactions, and simple clinical rules were enforced (e.g., ibuprofen excluded for hypertensive patients). Each patient record includes structured attributes: age (sampled uniformly from 1–90 years), binary sex encoding, and a random subset of comorbidity flags (diabetes, hypertension, asthma). The final dataset comprises approximately 50,000 samples with multi-label medication outputs.

Text preprocessing included lowercasing, removal of special characters, whitespace normalization, and drug alias standardization (e.g., *acetaminophen* $\rightarrow$ *paracetamol*). Disease name mismatches across sources were resolved via exact matching, substring matching, and fuzzy matching (via `difflib`) as a fallback. Non-clinical Synthea entries (e.g., "Housing unsatisfactory") were removed to retain only medically relevant conditions.

B. Models

Five neural architectures of increasing complexity were implemented and compared under identical training conditions to assess the contribution of different design choices to recommendation performance:

- **MLP:** A baseline model that flattens token embeddings and concatenates structured features before a linear output layer.
- **TextCNN:** A convolutional model applying 1D filters over token embeddings followed by adaptive max-pooling to capture local n-gram features.
- **BiLSTM:** A bidirectional LSTM that processes the token sequence and concatenates the final hidden states from both directions with structured features.
- **SelfAttention:** A single-head self-attention model that computes query–key–value attention over token embeddings and aggregates via mean pooling.
- **FusionNet:** A hybrid model combining a unidirectional LSTM for text encoding with a dedicated MLP branch for structured features, whose outputs are concatenated prior to the final classifier.

All models share a common vocabulary-based tokenizer with a maximum sequence length of 20 tokens and a 64-dimensional word embedding layer. Structured inputs (age, sex, three comorbidity flags) form a 5-dimensional feature vector. All models output logits over the full drug vocabulary via a linear layer, with Binary Cross-Entropy with Logits loss used for multi-label training.

C. Training Configuration

All models were trained using the Adam optimizer with a learning rate of 1×10^{-3} for 5 epochs with a batch size of 32. The dataset was split into 80% training and 20% test sets using a fixed random seed of 42. No validation set was held out during the comparative experiments; model selection is based on test-set evaluation after training convergence.

D. Evaluation Metrics

Model performance was assessed using a comprehensive set of multi-label classification metrics computed on the held-out test set. Threshold-based metrics were computed at a decision threshold of 0.5, while ranking metrics were computed at $k = 5$ by default. The following metrics were used:

- **Jaccard Index:** Ratio of the intersection to the union of predicted and ground-truth drug sets.
- **Micro / Macro F1, Precision, Recall:** Standard multi-label classification metrics aggregated at both the instance and class levels.
- **Hamming Loss:** Fraction of incorrectly predicted labels.
- **Subset Accuracy:** Fraction of samples where the predicted label set exactly matches the ground truth.
- **ROC-AUC (micro):** Area under the receiver operating characteristic curve aggregated across all drug classes.
- **PR-AUC:** Area under the precision–recall curve, computed using sample-averaged average precision.

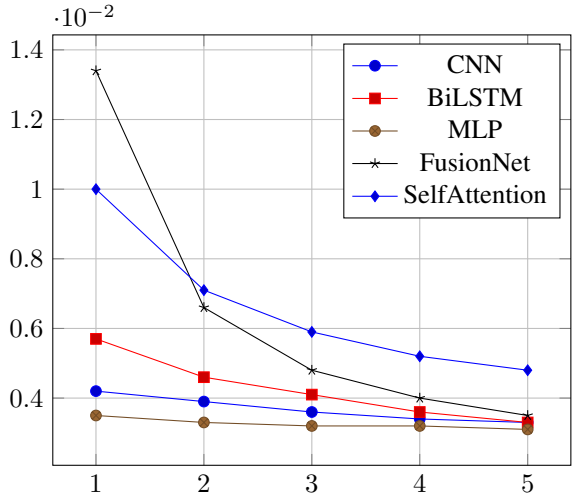

Fig. 2. Training Loss of various models across epochs.

- **Precision@K, Recall@K, F1@K:** Top-k ranking metrics evaluating the quality of the top-5 predicted medications per patient.
- **DDI Rate:** Fraction of predicted drug pairs in the top-k recommendations that correspond to known adverse interactions in the DDI dataset, used as a direct safety metric.

Top-K sensitivity was additionally evaluated for $k \in \{1, 3, 5, 10\}$ to assess ranking quality across different recommendation list lengths.

VI. Results

A. Training Convergence

Figure 2 shows the training loss curves for all five models across five epochs. The MLP exhibits the lowest absolute loss throughout training, beginning at 0.0035 and converging smoothly to 0.0031, which reflects the simplicity of its optimization landscape rather than superior generalization. The CNN and BiLSTM follow similar convergence trajectories, both reaching a final loss of approximately 0.0033 after starting at 0.0042 and 0.0057 respectively. FusionNet and SelfAttention begin with notably higher initial losses (0.0134 and 0.0100), indicating greater model complexity and a more challenging optimization surface, but both descend steeply in the first two epochs. FusionNet converges to 0.0035 by epoch four, while SelfAttention stabilizes at 0.0048, the highest final loss among all models. The faster loss reduction observed in FusionNet after the initial epochs suggests that its structured feature branch aids convergence once the text encoder begins to stabilize.

B. Multi-Label Classification Performance

Table II reports the full set of evaluation metrics across all models. FusionNet achieves the strongest overall performance, attaining the highest Jaccard index (0.7869), micro-precision (0.9252), micro-F1 (0.8808), and subset accuracy (0.6455), with a Hamming loss of 0.0010. BiLSTM performs comparably, ranking second on Jaccard (0.7796), micro-F1 (0.8762), and subset accuracy (0.6184). SelfAttention and CNN occupy an intermediate tier, with SelfAttention achieving marginally higher PRAUC (0.8764) and Precision@K (0.7114) than CNN and BiLSTM, suggesting that attention-based ranking is particularly effective for top-k retrieval. The MLP baseline, while competitive on micro-precision (0.8853), lags considerably on subset accuracy (0.4210) and Jaccard (0.6887), indicating that flat feature representations are insufficient for accurate multi-label drug set prediction.

Macro-averaged metrics are uniformly low across all models (F1_macro $\approx$ 0.07–0.08), which is expected given the large and imbalanced drug vocabulary; many drug classes appear rarely in the training data, resulting in poor per-class recall for infrequent medications.

C. Ranking and Top-K Performance

All models achieve high ROC-AUC scores ($>$ 0.993), confirming strong discriminative ability at the label level. PRAUC scores range from 0.8372 (MLP) to 0.8764 (SelfAttention), with CNN, BiLSTM, and FusionNet clustered closely around 0.874–0.875. At $k = 5$, Precision@K ranges from 0.6855 (MLP) to 0.7114 (SelfAttention), while F1@K spans 0.5825 to 0.6052. The narrow spread in top-K metrics among the four stronger models suggests that beyond a certain architectural

TABLE II
RESULTS OF THE EVALUATED MODELS ACROSS VARIOUS METRICS.

Model	**Jaccard**	**Precision_micro**	**Recall_micro**	**F1_micro**	**Precision_macro**	**Recall_macro**	**F1_macro**
MLP	0.6887	0.8853	0.7561	0.8156	0.0781	0.0675	0.0721
CNN	0.7657	0.9031	0.8342	0.8673	0.0858	0.0762	0.0798
BiLSTM	0.7796	0.9218	0.8348	0.8762	0.0899	0.0777	0.0821
SelfAttention	0.7684	0.9022	0.8382	0.8691	0.0864	0.0753	0.0789
FusionNet	0.7869	0.9252	0.8404	0.8808	0.0898	0.0786	0.0828
Model	**Hamming_Loss**	**Subset_Accuracy**	**ROC_AUC**	**PRAUC**	**Precision@K**	**Recall@K**	**F1@K**
MLP	0.0015	0.4210	0.9937	0.8372	0.6855	0.6329	0.5825
CNN	0.0011	0.5603	0.9968	0.8743	0.7101	0.6563	0.6042
BiLSTM	0.0010	0.6184	0.9957	0.8743	0.7076	0.6555	0.6028
SelfAttention	0.0011	0.5675	0.9970	0.8764	0.7114	0.6573	0.6052
FusionNet	0.0010	0.6455	0.9964	0.8753	0.7079	0.6561	0.6034

threshold, marginal gains in ranking quality diminish, and other factors such as safety filtering and interpretability become the primary differentiators.

D. Summary

FusionNet emerges as the best-performing architecture overall, combining strong multi-label accuracy with the highest subset accuracy, making it the most suitable backbone for the proposed medication recommendation framework. Its explicit separation of text and structured feature encoding aligns well with the multi-modal nature of outpatient clinical data. SelfAttention offers competitive ranking performance and may be preferred in scenarios where top-k precision is the primary concern. The MLP baseline, despite its simplicity, provides a reasonable lower bound and confirms that even shallow models benefit substantially from structured feature integration.

VII. DISCUSSION

The experimental results demonstrate that architectural complexity correlates positively with multi-label recommendation quality up to a point, beyond which gains plateau. FusionNet's design choice of maintaining separate encoding pathways for symptom text and structured patient features proves particularly well-suited to the outpatient setting, where demographic and comorbidity signals carry meaningful clinical weight alongside free-text complaints. The explicit fusion of these two modalities, rather than simple concatenation at the embedding level as in MLP and CNN, enables the model to learn richer cross-modal interactions, reflected in its superior subset accuracy and Jaccard index.

The consistently high ROC-AUC scores across all models suggest that the underlying symptom–medication signal in the synthetic dataset is learnable even by shallow architectures. However, the large gap in subset accuracy between MLP (0.4210) and FusionNet (0.6455) underscores that correctly predicting the full medication set for a patient, rather than individual drug labels in isolation, is a substantially harder problem that benefits from sequential or attention-based inductive biases. This distinction is clinically significant: a system that recommends most drugs correctly but omits one critical medication or includes one unsafe drug is of limited utility in practice.

The low macro-averaged F1 scores observed uniformly across all models highlight a known challenge in medication recommendation: the long-tail distribution of drug classes. Infrequent medications, which may be the most critical in edge cases such as rare comorbidities or drug allergies, are systematically under-predicted. This limitation is partly an artifact of the synthetic data generation process, where drug assignment is driven by disease–drug mappings rather than physician behavior, and may underrepresent clinically nuanced prescribing decisions. Addressing this imbalance through class-weighted loss functions, oversampling strategies, or prototype-based learning for rare drug classes is an important direction for future work.

SelfAttention's strong PRAUC and Precision@K performance relative to its subset accuracy suggests that attention mechanisms are particularly effective at surfacing the most relevant drugs at the top of the ranked list, even when the full predicted set is not exact. This property is desirable in a clinical decision support context, where a ranked shortlist presented to the clinician is often more actionable than a fixed predicted set. Integrating self-attention as a ranking re-scorer on top of FusionNet's predictions is a promising architectural extension.

A key limitation of the current evaluation is the absence of an explicit DDI rate comparison across models at inference time, as the safety filtering layer described in the methodology was applied during dataset construction rather than as a post-hoc

inference component in these experiments. Future work should decouple the safety filter and measure the raw DDI rate of model outputs before and after filtering, enabling a direct quantification of the safety benefit provided by the graph-based interaction layer. Additionally, all experiments were conducted on synthetically generated data, and performance on real-world clinical datasets may differ considerably due to the distributional gap between template-constructed symptom text and authentic clinical notes.

VIII. Conclusion

This paper presented a survey of twenty deep learning approaches to medication recommendation, identified persistent gaps in outpatient generalization, safety integration, and interpretability, and proposed a hybrid architecture addressing these limitations. Five neural models of increasing complexity were implemented and evaluated on a synthetic yet clinically grounded multi-label dataset constructed by integrating symptom–disease mappings, drug–disease associations, DDI constraints, and Synthea demographic records.

FusionNet, which combines a sequential text encoder with a dedicated structured feature branch, achieved the strongest overall performance across Jaccard index, micro-F1, and subset accuracy, demonstrating the value of modality-aware feature fusion for outpatient medication recommendation. SelfAttention showed competitive top-k ranking performance, suggesting its suitability in clinician-facing shortlist generation. All models achieved high ROC-AUC scores, confirming that the symptom-to-medication mapping signal is robustly learnable, while low macro-F1 scores highlighted the difficulty of rare drug prediction as an open challenge.

The proposed framework integrates a DDI-aware safety filter and SHAP-based interpretability layer to address clinical trust and patient safety requirements beyond raw predictive performance. Together, these components form a foundation for a lightweight, safety-aware, and explainable medication recommendation system suitable for primary-care decision support under the data-sparse conditions typical of outpatient settings.

Future work will focus on evaluation on real-world clinical datasets, incorporation of class-imbalance mitigation strategies for rare drug classes, extension of the safety filter to runtime inference with measurable DDI rate reduction, and exploration of larger pretrained clinical language models as drop-in symptom encoders within the FusionNet backbone.

References

[1] E. Choi *et al.*, "Doctor ai: Predicting clinical events via recurrent neural networks," in *MLHC*, 2016.

[2] E. Choi *et al.*, "Retain: An interpretable predictive model for healthcare using reverse time attention mechanism," in *NeurIPS*, 2016.

[3] J. Shang *et al.*, "Gamenet: Graph augmented memory networks for recommending medication combination," in *KDD*, 2019.

[4] W. Lee *et al.*, "Safedrug: Dual molecular graph encoders for safe drug recommendations," in *The Web Conference (WWW)*, 2021.

[5] X. Zhou, Y. Li, and W. Liang, "Cnn-rnn based intelligent recommendation for online medical pre-diagnosis support," *IEEE Trans. Emerg. Topics Comput.*, 2021.

[6] R. Miotto, L. Li, B. A. Kidd, and J. T. Dudley, "Deep patient: An unsupervised representation to predict the future of patients from the electronic health records," *Scientific Reports*, vol. 6, p. 26094, 2016.

[7] E. Choi *et al.*, "Gram: Graph-based attention model," in *KDD*, 2017.

[8] Y. Zhang *et al.*, "Leap: Learning to prescribe effective and safe treatment combinations for multimorbidity," in *KDD*, 2017.

[9] N. (various), "Molerec and molecular gnn approaches," *WWW/ICML Papers*, 2022–2023.

[10] E. Alsentzer *et al.*, "Publicly available clinical bert embeddings," in *NAACL ClinicalNLP*, 2019.

[11] J. Wang *et al.*, "Memory-augmented models for sequential recommendation," in *AAAI*, 2019.

[12] Z. Yang *et al.*, "Hierarchical attention networks for document classification," in *NAACL*, 2016.

[13] Z. Che *et al.*, "Recurrent neural networks for multivariate time series with missing values," *Scientific Reports*, 2018.

[14] M. Zitnik *et al.*, "Modeling polypharmacy side effects with graph convolutional networks," *Bioinformatics*, 2018.

[15] J. Lee *et al.*, "Transfer learning for clinical nlp," *Scientific Reports*, 2020.

[16] H. Harutyunyan *et al.*, "Multitask learning and benchmarking with clinical time series data," *Scientific Data*, 2019.

[17] M. Komorowski *et al.*, "The artificial intelligence clinician learns optimal treatment strategies for sepsis in intensive care," *Nature Medicine*, 2018.

[18] M. T. Ribeiro, S. Singh, and C. Guestrin, ""why should i trust you?": Explaining the predictions of any classifier," in *KDD*, 2016.

[19] (Representative), "Various works on knowledge graph embeddings for clinical entities," 2017–2022.

[20] K. Huang *et al.*, "Large clinical transformer models," 2022.